

Gene Regulatory Network Inference Enhanced by Transformer and LLM

Yue Wang

Irving Institute for Cancer Dynamics and
Department of Statistics, Columbia University

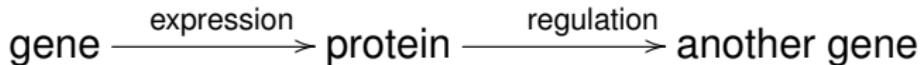
yw4241@columbia.edu

May 22, 2025

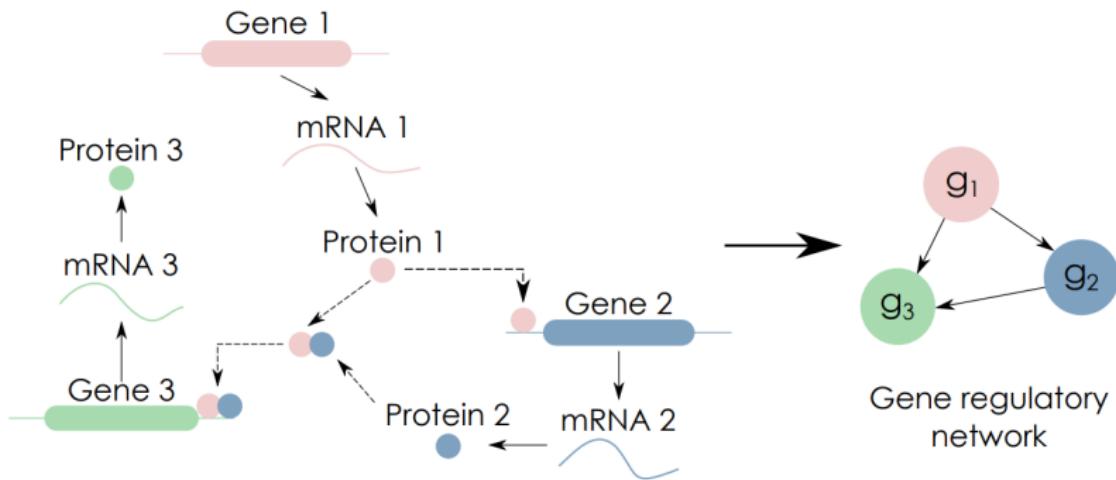
- Introduce gene regulatory network (GRN).
- Review the GRN inference method WENDY, and introduce its transformer-enhanced version, TRENDY.
- Introduce a novel GRN inference method, QWENDY, and its variants enhanced by transformers and large language models (LLMs).

Regulation of gene expression

- Gene expression: genes are transcribed to mRNAs and then translated to proteins.
- Various molecular regulators affect gene expression (change levels of mRNAs and/or proteins).
- Some regulators are small molecules, such as oxygen, sugars and vitamins. Some regulators are proteins. We focus on regulations between genes.

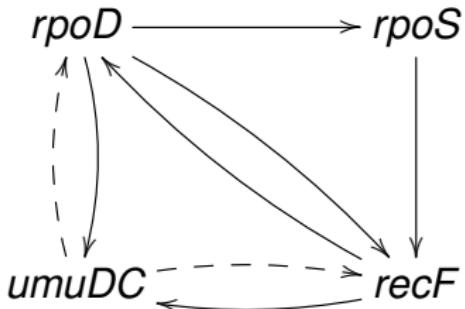


Regulation of gene expression



Genes and their regulatory relations form a gene regulatory network (GRN). It is a directed graph: vertices are genes, and edges are regulatory relations.

Regulation of gene expression



- An example of GRN in *E. coli*. Each vertex is a gene. Two types of regulations: solid arrow means activation, and dashed arrow means inhibition.
- GRN determines cell function. If the GRN structure is known, we know how the cell fate is determined, and how we can change the cell fate.

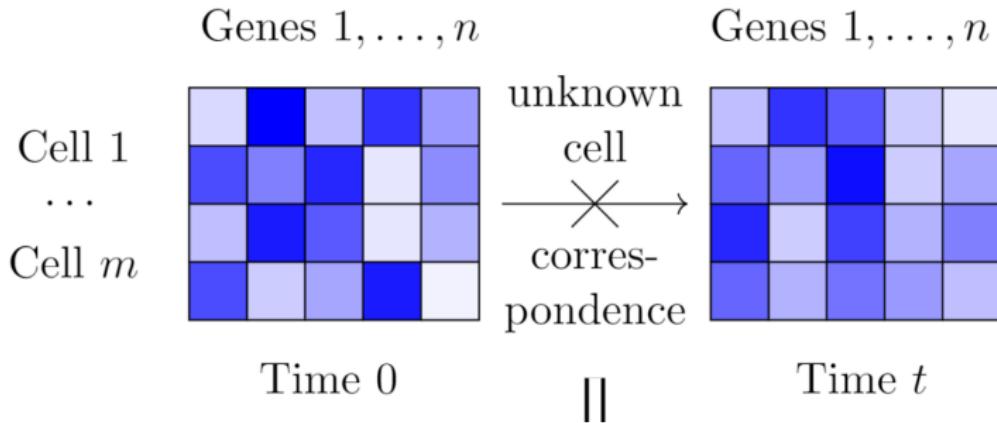
Regulation of gene expression

- A central question in biology is to determine the GRN structure.
- For two genes V_i, V_j , does the expression of V_i regulate (activate or inhibit) the expression of V_j ?
- Genes (DNAs), mRNAs and proteins are generally confined within living cells.
- It is extremely difficult or even impossible to directly determine whether one gene regulates another gene with biochemical methods.
- We have accumulated a large amount of gene expression data. Certain types of gene expression data can be used to infer the GRN structure.

Data type

- We consider a specific data type, which is common in experiment and very informative, but there are very few GRN inference methods developed for it specifically.
- For each cell, we can use single-cell RNA sequencing to measure the expression levels (mRNA counts) of different genes.
- Since single-cell level gene expression is stochastic, we repeat this for different cells. The data are represented as a matrix with m rows (cell number) and n columns (gene number), namely m samples of an n -dimensional random variable.
- We do this at multiple time points. If the cells are not at stationary, different time points have different distributions and more information.
- At different time points, we measure different cells, since measurement kills cells.

single-cell gene expression data



At time 0 and t , we obtain expression data $X(0)$ and $X(t)$, where each has m samples for an n -dimensional random variable.

- In a previous paper, we developed the WENDY method for this type of data.
- We can obtain a simplified relation for $X(0)$ and $X(t)$:

$$X(t) \approx X(0)(I + tA) + tb + X(0) \odot \epsilon(t).$$

- Here A is an $n \times n$ matrix, representing the GRN we want: $A_{i,j} > 0$ means gene i activates gene j ; $A_{i,j} < 0$ means gene i inhibits gene j ; and $A_{i,j} = 0$ means gene i does not regulate gene j directly.
- Besides, b is a $1 \times n$ vector, representing the base synthesis rate, $\epsilon(t)$ is random noise, and \odot is the entrywise (Hadamard) product.

- For n -dimensional $X(0)$ or $X(t)$, we can calculate the $n \times n$ covariance matrix K_0 or K_t .
- When sample number m is small, we can use graphical lasso to avoid K_0 being degenerate.
- We can obtain the dynamics of K_t :

$$K_t \approx (I + tA^T)K_0(I + tA).$$

- Now we have an optimization problem:

$$\min_A f(A) := \|K_t - (I + tA^T)K_0(I + tA)\|_F^2.$$

- This non-convex optimization problem can be solved numerically with BFGS or other methods.
- WENDY method outputs the GRN A from K_0 and K_t .

- In a recent paper, we developed a method to enhance WENDY.
- In practice, K_t does not accurately equal to $(I + tA^T)K_0(I + tA)$.
- Therefore, even if the numerical solver works perfectly, we might not obtain the accurate GRN A .
- Define

$$K_t^* = (I + tA^T)K_0(I + tA).$$

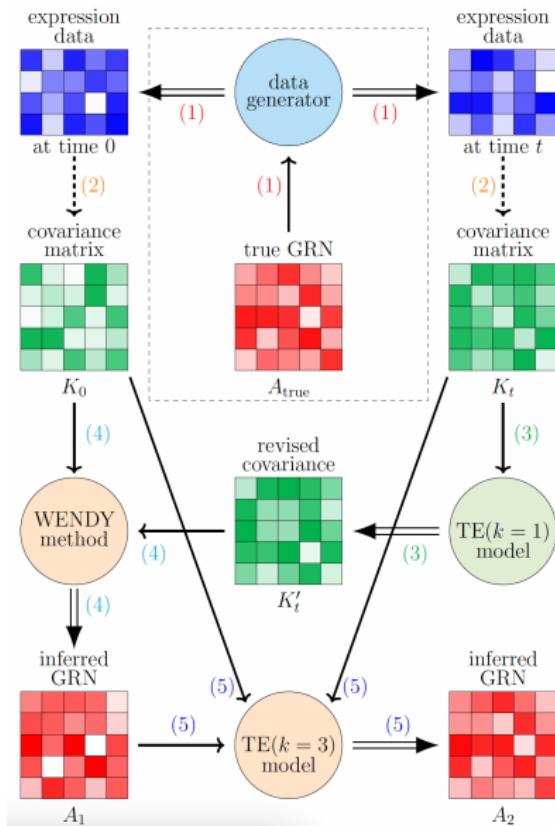
- If we feed K_0 and K_t^* to the solver, we could get a better A .
- We know K_0 and K_t from data, but A is unknown, so we cannot calculate K_t^* .

- We can use some simulator to generate many gene expression data with random ground truth GRN A_{true} :

$$dX_j(t) = V \left\{ \beta \prod_{i=1}^n \left[1 + (A_{\text{true}})_{i,j} \frac{X_i(t)}{X_i(t) + 1} \right] - \theta X_j(t) \right\} dt + \sigma X_j(t) dW_j(t).$$

- Then train a supervised-learning model with input K_t and target K_t^* (which is known for training data), so the output K'_t is close to K_t^* .
- In reality, when we only have K_0 and K_t , use this model to output K'_t , and feed K_0 and K'_t to the solver to obtain a better GRN, A_1 .
- We can train another supervised-learning model with inputs A_1 , K_0 , K_t and target A_{true} . The output A_2 should be better than A_1 .

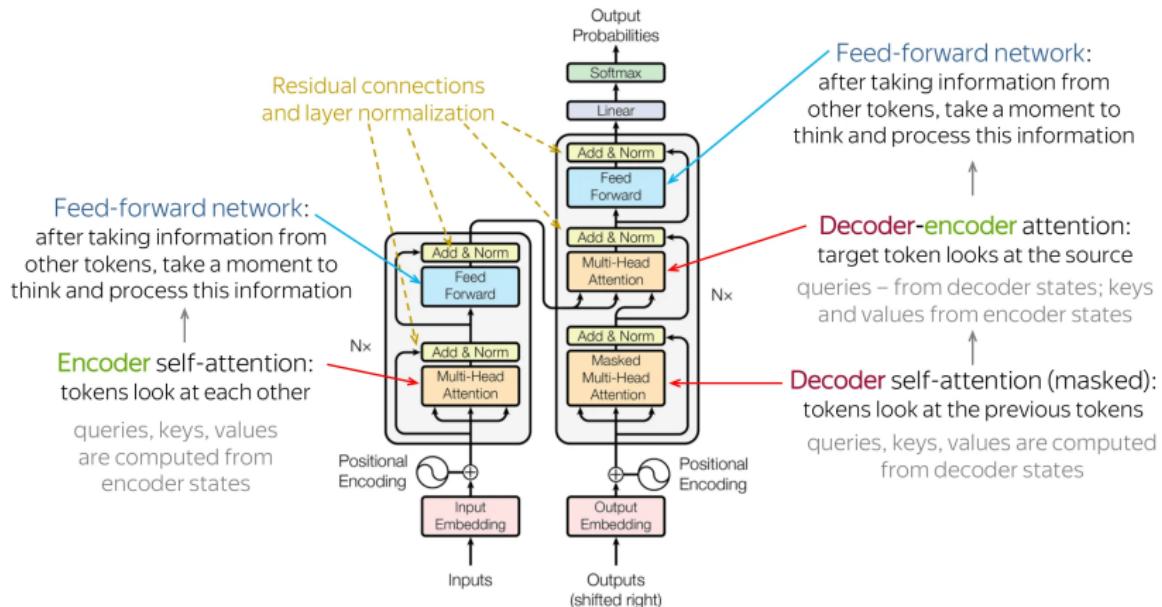
TRENDY method



- What supervised-learning model can be used for this task?
- The input is one or several matrices, and the target is a matrix. It is a highly nonlinear regression problem, and simple linear/nonlinear regression does not work.
- Since we have a data simulator, which leads to a large amount of data, we can train a deep learning model with many parameters.
- In TRENDY, the core of this model is the transformer encoder.

- The transformer model, introduced by Vaswani et al. in 2017, is a neural network architecture designed for sequence modeling tasks such as translation and text generation.
- The transformer architecture has become the foundation of many modern models, including BERT, GPT, and T5.
- A transformer is composed of an encoder and a decoder:
 - The **encoder** consists of multiple identical layers that process the input sequence and generate context-rich embeddings.
 - The **decoder** also consists of multiple layers and generates the output sequence by attending to both previous decoder outputs and the encoder output.

TRENDY method



- For our highly nonlinear regression task, we can use the transformer encoder.
- Transformer encoder can be regarded as an oracle machine that learns the features of the input, which work well for downstream tasks.
- For transformer encoder, the input is m tokens, each as a d_{model} -dimensional vector, thus being a matrix of size $m \times d_{\text{model}}$. The output is also an $m \times d_{\text{model}}$ matrix. Here d_{model} is called the model dimension.
- (Optional): Details of transformer encoder and TRENDY.

- The core component of transformer encoder, where X is the $m \times d_{\text{model}}$ input:

$$Q = XW^Q, \quad K = XW^K, \quad V = XW^V,$$

$$\text{Attention} = \text{softmax} \left(\frac{QK^T}{\sqrt{d_{\text{model}}}} \right) V.$$

Here W^Q, W^K, W^V (query, key, value) are $d_{\text{model}} \times d_{\text{model}}$ trainable matrices.

- It is the self-attention mechanism, which allows each token in a sequence to attend to all other tokens, capturing long-range dependencies effectively.

1. **Input:** k matrices (k groups of $n \times n$)
2. Linear embedding layer with dimension d (k groups of $n \times n \times d$)
3. Segment embedding layer (omitted when $k = 1$) (k groups of $n \times n \times d$)
4. 2-D positional encoding layer (k groups of $n \times n \times d$)
5. Flattening and concatenation ($(n^2) \times (dk)$)
6. l layers of transformer encoder ($(n^2) \times (dk)$)
7. Linear embedding layer ($n^2 \times 1$)
8. **Output:** reshaping into a matrix ($n \times n$)

Structure of the $\text{TE}(k)$ model used in TRENDY

- Since the transformer encoder does not have a built-in notion of sequence order, **positional encodings** are added to input embeddings before they are fed into the model, so the position information of each token is known to the model.

$$\text{PE}_{\text{pos},2i} = \sin \left(\frac{\text{pos}}{10000^{\frac{2i}{d_{\text{model}}}}} \right)$$

$$\text{PE}_{\text{pos},2i+1} = \cos \left(\frac{\text{pos}}{10000^{\frac{2i}{d_{\text{model}}}}} \right)$$

Since our input is 2-D matrices, we need a 2-D positional encoding.

$$\text{PE}[x, y, 2j - 1] = \text{PE}[x, y, 2j - 1 + d/2] = \sin[(y - 1) \times 10^{-32(j-1)/d}],$$

$$\text{PE}[x, y, 2j] = \text{PE}[x, y, 2j + d/2] = \cos[(y - 1) \times 10^{-32(j-1)/d}],$$

$$\text{PE}[x, y, 2j - 1 + d/4] = \text{PE}[x, y, 2j - 1 + 3d/4] = \sin[(x - 1) \times 10^{-32(j-1)/d}],$$

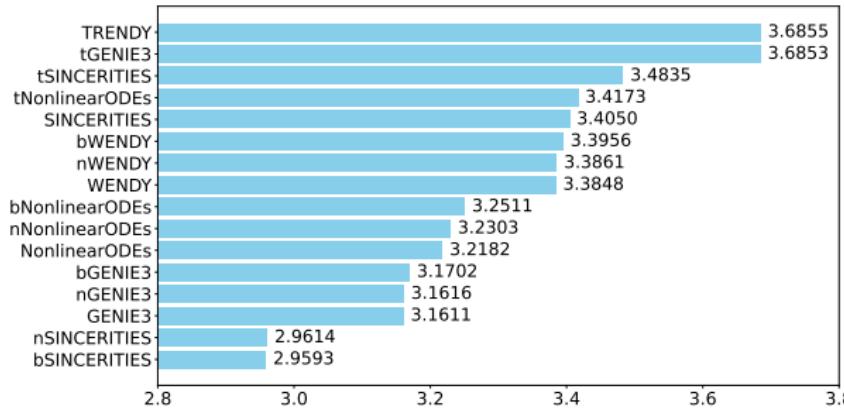
$$\text{PE}[x, y, 2j + d/4] = \text{PE}[x, y, 2j + 3d/4] = \cos[(x - 1) \times 10^{-32(j-1)/d}].$$

- Since the model does not distinguish between different input matrices, we need to add different **segment embedding** vectors to different input matrices.
- The segment embedding vector is trainable with length d_{model} , and it is the same for all tokens in the same input matrix.

- WENDY method uses two covariance matrices to infer the GRN.
- WENDY relies on a linear approximation, which leads to an error.
- TRENDY uses the transformer encoder to learn this error and decrease it.
- This idea can be used to enhance other GRN inference methods: use the inferred GRN as input and true GRN as target, and train a similar transformer encoder.
- There are other methods to enhance GRN inference, but they cannot determine the direction of regulation.

TRENDY method

- We test 16 methods: four base methods, their transformer-enhanced versions and two other enhanced version.
- There are two synthetic data sets, DREAM4 and SINC, and two experimental data sets, hESC and THP-1.
- We use two quantities to measure the performance: AUROC and AUPRC. They are between 0 (perfect mismatch) and 1 (perfect match).
- Here is the sum of AUROC and AUPRC on four data sets.



- We introduce TRENDY, which uses the transformer encoder to enhance WENDY.
- With covariance matrices from two time points, WENDY solves a non-convex optimization problem with infinitely many solutions.
- With covariance matrices from more time points, can we infer the GRN easier?
- With data from four time points, we can uniquely determine the GRN just by matrix factorization.
- This new method is called QWENDY, where Q means quadruple.

- We have data at time points $0, t, 2t, 3t$, each as m samples for n genes. Their $n \times n$ covariance matrices are K_0, K_1, K_2, K_3 . For the GRN matrix A , define $B = I + tA$.
- We have approximated relations:

$$K_1 = B^T K_0 B, \quad K_2 = B^T K_1 B, \quad K_3 = B^T K_2 B. \quad (1)$$

- Define the $1 \times n$ expected levels at four time points as $\mathbf{x}_0, \mathbf{x}_1, \mathbf{x}_2, \mathbf{x}_3$. They have approximated relations:

$$\mathbf{x}_1 = \mathbf{x}_0 B + t\mathbf{c}, \quad \mathbf{x}_2 = \mathbf{x}_1 B + t\mathbf{c}, \quad \mathbf{x}_3 = \mathbf{x}_2 B + t\mathbf{c}, \quad (2)$$

where \mathbf{c} is an unknown vector.

QWENDY method

- Given $K_0, K_1, K_2, K_3, \mathbf{x}_0, \mathbf{x}_1, \mathbf{x}_2, \mathbf{x}_3$, we will solve B from Eq. 1 and Eq. 2.
- Step (1):** For any B , we want to minimize

$$\|K_1 - B^T K_0 B\|_F^2. \quad (3)$$

- Consider Cholesky decompositions

$$K_1 = L_1 L_1^T, \quad K_0 = L_0 L_0^T,$$

where L_1 and L_0 are lower-triangular and invertible.

- Define

$$O = L_1^{-1} B^T L_0, \quad B = L_0^{-T} O^T L_1^T,$$

then the target Eq. 3 becomes

$$\|L_1 L_1^T - L_1 O O^T L_1^T\|_F^2.$$

Therefore, Eq. 3 is minimized to 0 if and only if O is orthonormal: $O O^T = I$.

- We use K_0 and K_1 to restrict B to a space with the same dimension as the set of all orthonormal matrices.

- **Step (2):** For such B that minimizes Eq. 3, we want to find B that makes $B^T K_1 B$ close to K_2 .
- Here we do not minimize

$$\|K_2 - B^T K_1 B\|_F^2$$

as it is difficult. Instead, we want to minimize

$$\|L_1^{-1}(K_2 - B^T K_1 B)L_1^{-T}\|_F^2 \quad (4)$$

among B that minimizes Eq. 3.

- Consider the eigenvalue decomposition

$$L_0^{-1} K_1 L_0^{-T} = P_1 D_1 P_1^T,$$

where P_1 is orthonormal, and D_1 is diagonal with strictly increasing positive diagonal elements (eigenvalues)

- We also have

$$L_1^{-1} K_2 L_1^{-T} = P_2 D_2 P_2^T$$

with orthonormal P_2 and diagonal D_2 with strictly increasing positive diagonal elements (eigenvalues).

- Now the target Eq. 4 equals

$$\|P_2 D_2 P_2^T - O P_1 D_1 P_1^T O^T\|_F^2 = \|D_2 - P_2^T O P_1 D_1 P_1^T O^T P_2\|_F^2,$$

since P_2 is orthonormal and does not affect Frobenius norm.

- Define

$$W = P_2^T O P_1, \quad O = P_2 W P_1^T, \quad B = L_0^{-T} P_1 W P_2^T L_1^T,$$

where W is orthonormal.

- Then Eq. 4 equals

$$\|D_2 - WD_1 W^T\|_F^2. \quad (5)$$

- We use the following lemma to handle Eq. 5:

Lemma

For diagonal D_1, D_2 with strictly increasing diagonal elements and any orthonormal W , Eq. 5 is minimized when W is diagonal, and the diagonal elements are ± 1 . (There are 2^n possibilities for such W .)

- Given K_0, K_1, K_2 , we can restrict B to 2^n possibilities.

- **Step (3):** For such B that minimizes Eq. 4 among B that minimizes Eq. 3, we want to find B that makes $B^T K_2 B$ close to K_3 .
- Here we do not minimize

$$\|K_3 - B^T K_2 B\|_F^2$$

as it is difficult. Instead, we want to minimize

$$\|L_1^{-1}(K_3 - B^T K_2 B)L_1^{-T}\|_F^2. \quad (6)$$

- Define

$$G = P_2^T L_1^{-1} K_3 L_1^{-T} P_2,$$

and

$$H = P_1^T L_0^{-1} K_2 L_0^{-T} P_1.$$

- Eq. 6 equals

$$\begin{aligned}
 & \|P_2^T L_1^{-1} K_3 L_1^{-T} P_2 - W P_1^T L_0^{-1} K_2 L_0^{-T} P_1 W\|_F^2 \\
 &= \|G - WHW\|_F^2 = \|G\|_F^2 + \|W H W\|_F^2 - 2\|G \otimes (W H W)\|_F^2 \\
 &= \|G\|_F^2 + \|H\|_F^2 - 2\|G \otimes (W H W)\|_F^2,
 \end{aligned} \tag{7}$$

where \otimes is element-wise product.

- We want to maximize $\|G \otimes (W H W)\|_F^2$. Define $C = G \otimes H$, which is still positive definite and symmetric.
- Denote the diagonal elements of W by $\mathbf{w} = [w_1, \dots, w_n]$. Then we need to maximize

$$\|G \otimes (W H W)\|_F^2 = \mathbf{w} C \mathbf{w}^T.$$

- With general vector \mathbf{v} with $||\mathbf{v}||_2^2 = n$, $\mathbf{v}C\mathbf{v}^T$ is maximized by the eigenvector that corresponds to the largest eigenvalue of C (unique up to a \pm sign).
- After obtaining \mathbf{v} , we just need to project it to $[w_1, \dots, w_n]$ with $w_i = \pm 1$ by taking the sign of each term: $w_i = \text{sign}(v_i)$.
- Then we have W , and finally

$$B = L_0^{-T} P_1 W P_2^T L_1^T.$$

- Given K_0, K_1, K_2, K_3 , we can uniquely determine B up to a \pm sign.
- Since $B^T K B = (-B)^T K (-B)$, more K_i cannot provide more information.

- **Step (4):** For B and $-B$ from Step (3), we can use $\mathbf{x}_0, \mathbf{x}_1, \mathbf{x}_2, \mathbf{x}_3$ to determine which is better, since they cannot both satisfy Eq. 2.
- We just need to compare the errors of B and $-B$ and choose the smaller one as the output.

QWENDY method

1. **Input** expression levels of n genes over m cells at time points 0, t , $2t$, $3t$
2. **Calculate** covariance matrices K_0, K_1, K_2, K_3 , and mean levels $\mathbf{x}_0, \mathbf{x}_1, \mathbf{x}_2, \mathbf{x}_3$
3. **Calculate** Cholesky decomposition

$$K_1 = L_1 L_1^T, \quad K_0 = L_0 L_0^T$$

Calculate eigenvalue decomposition

$$L_0^{-1} K_1 L_0^{-T} = P_1 D_1 P_1^T, \quad L_1^{-1} K_2 L_1^{-T} = P_2 D_2 P_2^T$$

Calculate

$$C = (P_2^T L_1^{-1} K_3 L_1^{-T} P_2) \otimes (P_1^T L_0^{-1} K_2 L_0^{-T} P_1)$$

4. **Calculate** \mathbf{v} , the eigenvector that corresponds to the largest eigenvalue of C , and the projection \mathbf{w} with $w_i = \text{sign}(v_i)$
Construct W with \mathbf{w} on diagonal
5. **Calculate** $B = L_0^{-T} P_1 W P_2^T L_1^T$
Compare total squared errors of B and $-B$ for Eqs. 4–6
6. **Output** B or $-B$, the one that corresponds to the smaller error, and the GRN $A = (B - I)/t$

Workflow of QWENDY method.

- Since Eq. 1 and Eq. 2 are approximated, there are two problems:
- (1) if these equations hold accurately, can we solve B ;
- (2) if these equations do not quite hold, can we find B to minimize the error.
- The derivation of QWENDY solves problem (2). The following theorem shows that QWENDY also solves problem (1).

Theorem

If covariance matrices K_0, K_1, K_2, K_3 and mean levels $\mathbf{x}_0, \mathbf{x}_1, \mathbf{x}_2, \mathbf{x}_3$ satisfy Eq. 1 and Eq. 2 for some B_0 , then QWENDY will output B_0 .

- Similar to TRENDY, we can try to enhance QWENDY.
- One idea is to train a transformer encoder to calculate K'_1, K'_2, K'_3 that better fit Eq. 1. The corresponding enhanced method with 3M trainable parameters is named TEQWENDY.
- We can also try large language models (LLMs). In general, an LLM has three sections:
 - (1) Convert text to vector representations;
 - (2) Use transformer (encoder layers, decoder layers, or both) to learn contextual relationships;
 - (3) Convert model outputs back to natural language.
- LLMs are trained with massive natural language data.
- We can choose an LLM with pre-trained encoder layers and fine-tune it.

- We use the RoBERTa Large model with 24 transformer encoder layers and 300M parameters.
- Training all parameters is not feasible. We use LoRA method to fine-tune it.
- If the original pre-trained weight matrix W (frozen) has size $p \times q$, LoRA trains two low rank matrices A and B, with size $p \times r$ and $r \times q$, where $r \ll p, q$. The final weight matrix is $W + AB$.
- In our case, the number of trainable parameters decreases from 300M to 3M.
- We use the fine-tuned RoBERTa Large model to replace the transformer part in TEQWENDY, and name it LEQWENDY.

- The dimension of RoBERTa Large model is fixed (1024). We need to process the input to match this dimension.
- For both TEQWENDY and LEQWENDY, the training takes a few hours on a single PC with a good graphics card.
- The most time-consuming step is to generate enough training data (about two weeks).

QWENDY method

- Since TRENDY ranks the first in 16 methods, we only test new methods against TRENDY.
- We use the same data sets and measurements.
- TEQWENDY is better than TRENDY. LEQWENDY totally fails, although with the same amount of trainable parameters with TEQWENDY.
- QWENDY does not perform well on synthetic data, but it is the best in all 19 methods on two experimental data sets.

		QWENDY	LEQ-WENDY	TEQ-WENDY	TRENDY
SINC	AUROC	0.5314	0.4972	0.9587	0.8941
	AUPRC	0.5720	0.5185	0.9142	0.8314
DREAM4	AUROC	0.4987	0.5164	0.5372	0.5341
	AUPRC	0.1844	0.1823	0.2203	0.2177
THP-1	AUROC	0.5524	0.5543	0.5415	0.5557
	AUPRC	0.4294	0.3632	0.3801	0.3669
hESC	AUROC	0.6019	0.5905	0.4815	0.5311
	AUPRC	0.0435	0.0367	0.0317	0.0376
Total		3.4137	3.2591	4.0652	3.9686

- Review WENDY method and introduce its transformer-enhanced version, TRENDY.
- Introduce a novel GRN inference method, QWENDY, and its variants enhanced by transformers and large language models (LLMs).
- Transformer encoder is a powerful tool for general tasks. LLMs are trained with natural language, not very helpful to science. We should train different large foundation models for different subjects and fine-tune on specific tasks.

References

- Xueying Tian, Yash Patel, and Yue Wang. (2024). "TRENDY: gene regulatory network inference enhanced by transformer." bioRxiv, 2024.10.14.618189, accepted by Bioinformatics.
- Yue Wang, and Xueying Tian. (2025). "QWENDY: gene regulatory network inference enhanced by large language model and transformer." bioRxiv, 2025.02.22.639640.



Code files: <https://github.com/YueWangMathbio/TRENDY>
<https://github.com/YueWangMathbio/QWENDY>